


Introduction


Machine Learning with Big Data

Evert Duipmans
Jeroen Linssen
Etto Salomons



Contents

- Evaluation of Q1 of BDT: DBDP & IML
- Organization
- Recommendation systems
 - Content-based recommendations
 - Text learning
 - Cosine similarity



Evaluation of Q1 of BDT

Distributed Big Data Processing

tinyurl.com/BDT22dbdp

Introduction to Machine Learning

tinyurl.com/BDT22iml

Organization

Every week one lab session

- Some theory
- Lot of time to work on your exercises
- Rooms reserved in Deventer (X2.03) and Enschede (RB4.03)

Learning materials

- Slides
- Videos
- Articles
- Jupyter Notebooks



Organization: assignments

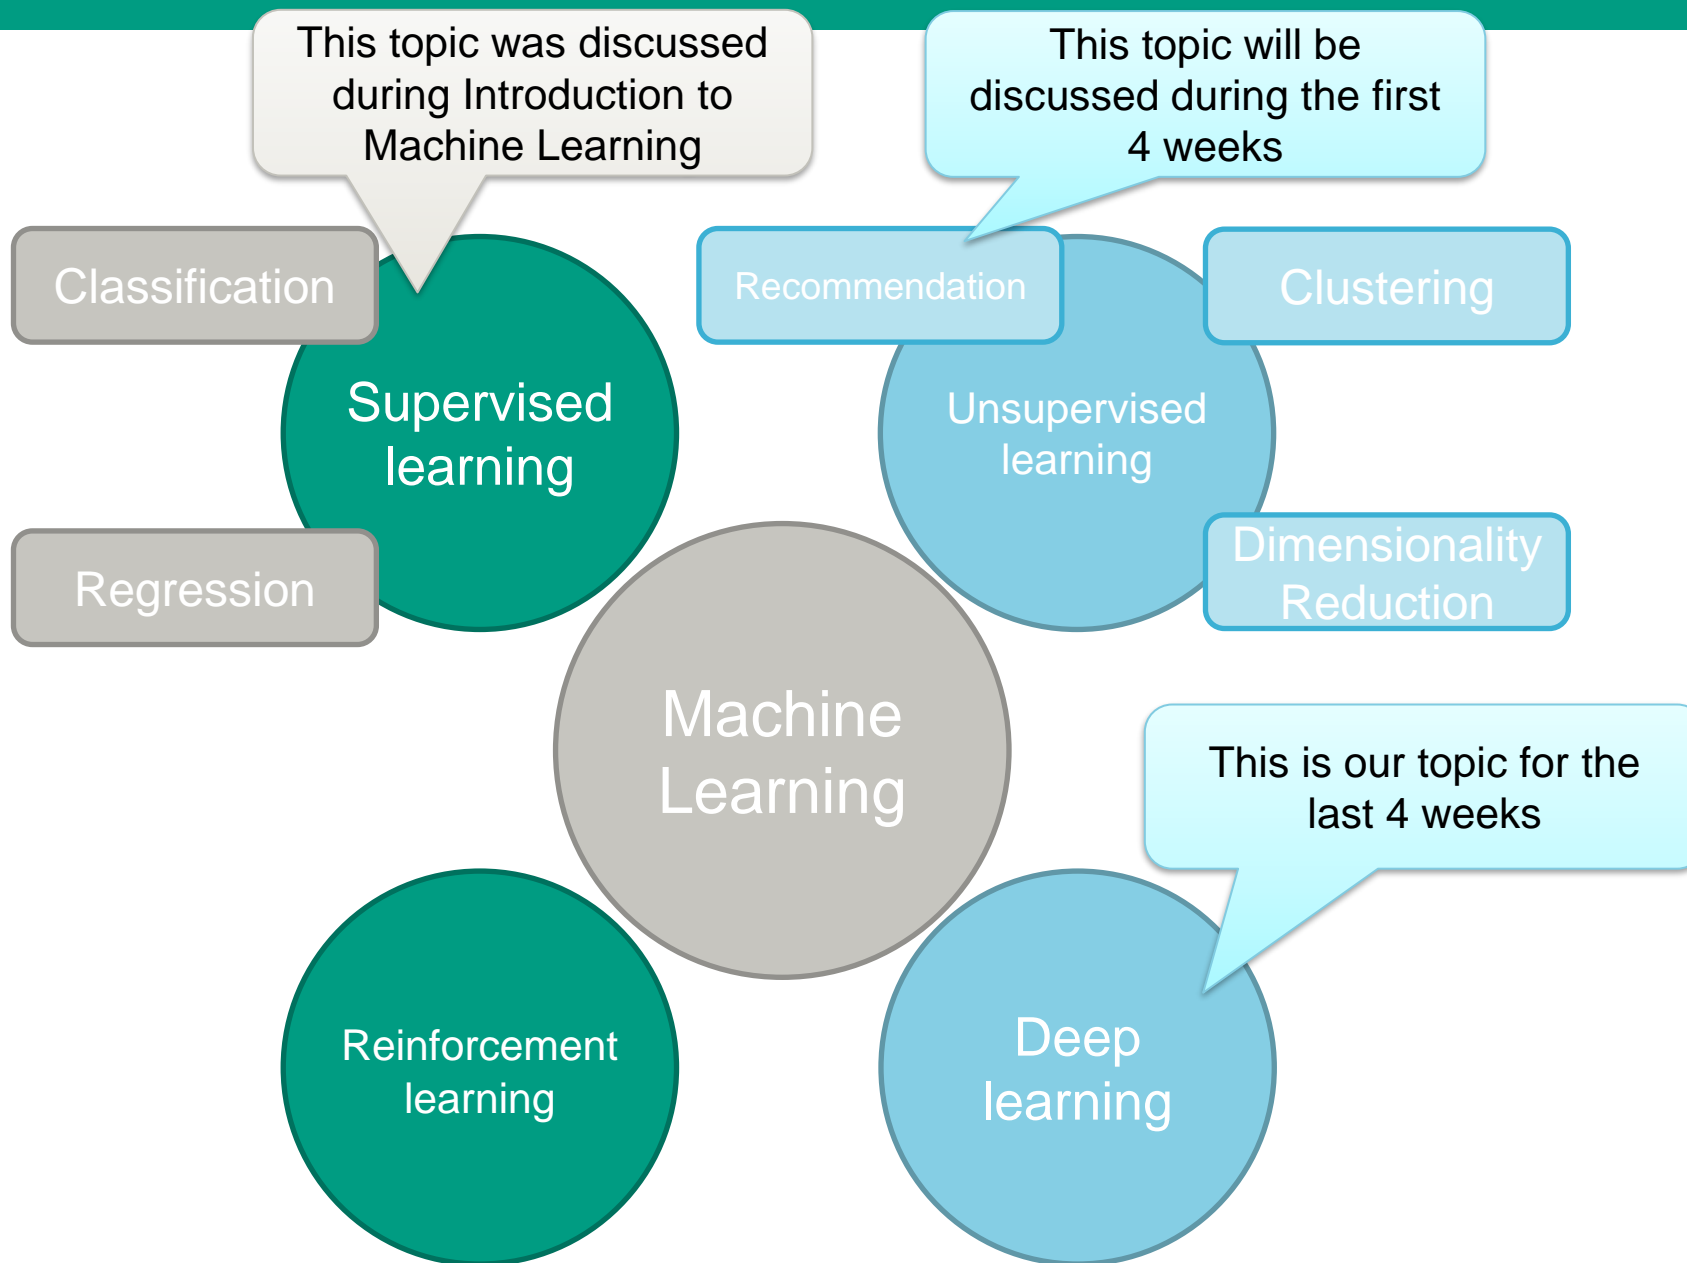
1. Recommender system (4 weeks)

Build your own recommendation system based on three different techniques.

2. Deep learning (4 weeks)

- Training neural networks using fully connected layers and convolutional layers
- Retraining existing neural networks

Grade = avg(assignment1, assignment2) if both ≥ 5

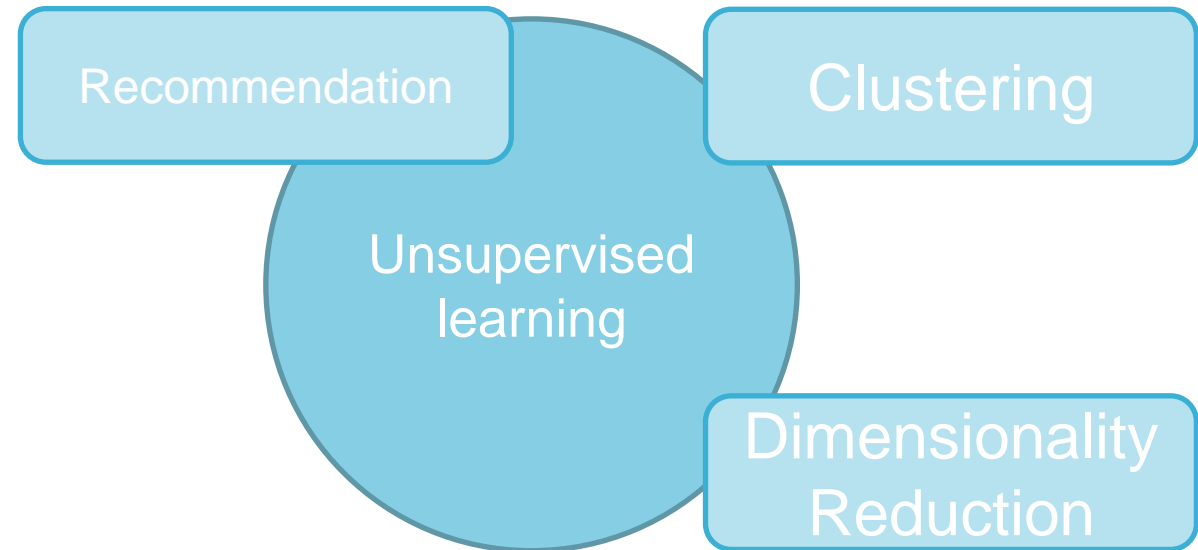


Unsupervised learning

Trains a **model** with "unlabeled" data.

Used for:


- Finding patterns in data
- Finding similar users
- Detecting anomalies



Content-based recommendations

Machine Learning with Big Data

Evert Duipmans
Jeroen Linssen
Etto Salomons





Coldplay - Paradise (Official Video)

1.043.306.530 weergaven



4,2 MLN.



129K



DELEN



OPSLAAN



Coldplay

Gepubliceerd op 18 okt. 2011

ABONNEREN 13 MLN.

Taken from the album Mylo Xyloto. Stream / download at <http://smarturl.it/cpmyloxloto>

~ Follow Coldplay ~

MEER WEERGEVEN



Ed Sheeran - Perfect (Official...)

Ed Sheeran

1,5 mld. weergaven

4:40



Eiffel 65 - Blue (Da Ba Dee)

Km Music

89 mln. weergaven

5:08



twenty one pilots: Stressed Out...

Fueled By Ramen

1,5 mld. weergaven

3:46



Tape Face Auditions & Performances |...

Got Talent Global

Aanbevolen voor jou

17:44



John Newman - Love Me Again

John Newman

621 mln. weergaven

3:56



OneRepublic - Counting Stars

OneRepublic

2,5 mld. weergaven

4:44



Coldplay's FULL Pepsi Super Bowl



bakery



Alle

Maps

Afbeeldingen

Shopping

Nieuws

Meer

Instellingen

Tools

Ongeveer 379.000.000 resultaten (0,54 seconden)



Beoordeling ▾ Openingstijden ▾

Specialist Jagers

3,9 ★★★★★ (24) · Bakkerij

Korte Hengelsestraat 20



Nienke's Cupcakes

4,9 ★★★★★ (57) · Gebak

Korte Haaksbergerstraat 15

Gesloten · Opent om 10:30




Nienke's Cupcakes

4,9 ★★★★★ (57) · Gebak

Korte Haaksbergerstraat 15

Gesloten · Opent om 10:30



Banketbakkerij Oonk Enschede

4,3 ★★★★★ (44) · Bakkerij

Kuipersdijk 21



☰ [Meer plaatsen](#)

Laura's Bakery - Van bakken tot borrelen

<https://www.laurasbakery.nl/> ▼

Laura's Bakery staat vol met de lekkerste toegankelijke recepten. Van hartig tot zoet en van bakken tot borrelen. Voor ieder wat wils!

Bakery - Wikipedia

<https://en.wikipedia.org/wiki/Bakery> ▼ [Vertaal deze pagina](#)

A **bakery** is an establishment that produces and sells flour-based food baked in an oven such as bread, cookies, cakes, pastries, and pies. Some retail **bakeries** ...

[History](#) · [Specialities](#) · [Commercialization](#)

Appèl: "Met The Bakery ontstijgen we de standaard catering"

www.schoolfacilities.nl/.../3736-appel-met-the-bakery-ontstijgen-we-de-standaard-cat... ▼

19 jan. 2016 - Zeer opvallend aanwezig is bijvoorbeeld The Bakery, één van de nieuwe concepten die Appèl bij Saxion heeft geïntroduceerd. De bakkerij ...

Recommender systems

What data is used?

Explicit ratings

- Rate content (stars, like/dislike, ...)
- Requires extra work for the user
- Cultural differences
- People rate different
- Data is often sparse



Implicit ratings

- Things you do: click on links, read article, add to cart, buy things, how long did you watch a video?
- Lots of companies use sales data
- Things you consume

Recommender systems

Top-N list

The overall goal of a recommendation systems is to:

Recommend n relevant items to the user



Recommender systems

Problems

1. Cold-start problem
 - First-time user
 - Not enough user interactions for a particular item
2. Sparsity problem
 - Users only rated a very small number of items
3. First-rater problem
 - Items that have not been rated will never be recommended

Recommender systems


Approaches

Two possible approaches to building recommender systems:

1. **Content-based (week 1)**
Recommend items with the same properties
2. **Collaborative filtering (week 2 and 3)**
Recommend based on ratings of similar users



Content-based recommendations



Content-based recommendations


General idea:

Recommend items to the user that are similar to items that the user already likes.

Examples

- Recommend books from the same authors or the same genre
- Recommend music from the same artist, collaborations or music with the same BPM

CONTENT-BASED FILTERING



Content-based recommendations

Music example

1. Create an **item profile** for each album in the catalog.
Each profile consists of features, such as: artists, titles, genre, band members, ...
2. Create a **user profile** consisting of the item profiles of the items that have been purchased by the user.
3. Find items that are similar to the items in the user profile.
4. Filter the list and recommend the top-N items.

Content-based recommendations


1. Create item profiles



Title	Artist	Members	Genre
Nevermind	Nirvana	Kurt Cobain	Rock, Grunge
Night visions	Imagine Dragons	Dan Reynolds	Rock
Hot fuss	Killers	Brandon Flowers	Rock, Pop
Greatest hits	Nickelback	Chad Kroeger	Awful
Vultures	Kensington	Eloi Youssef	Rock
Overexposed	Maroon 5	Adam Levine	Pop

Content-based recommendations

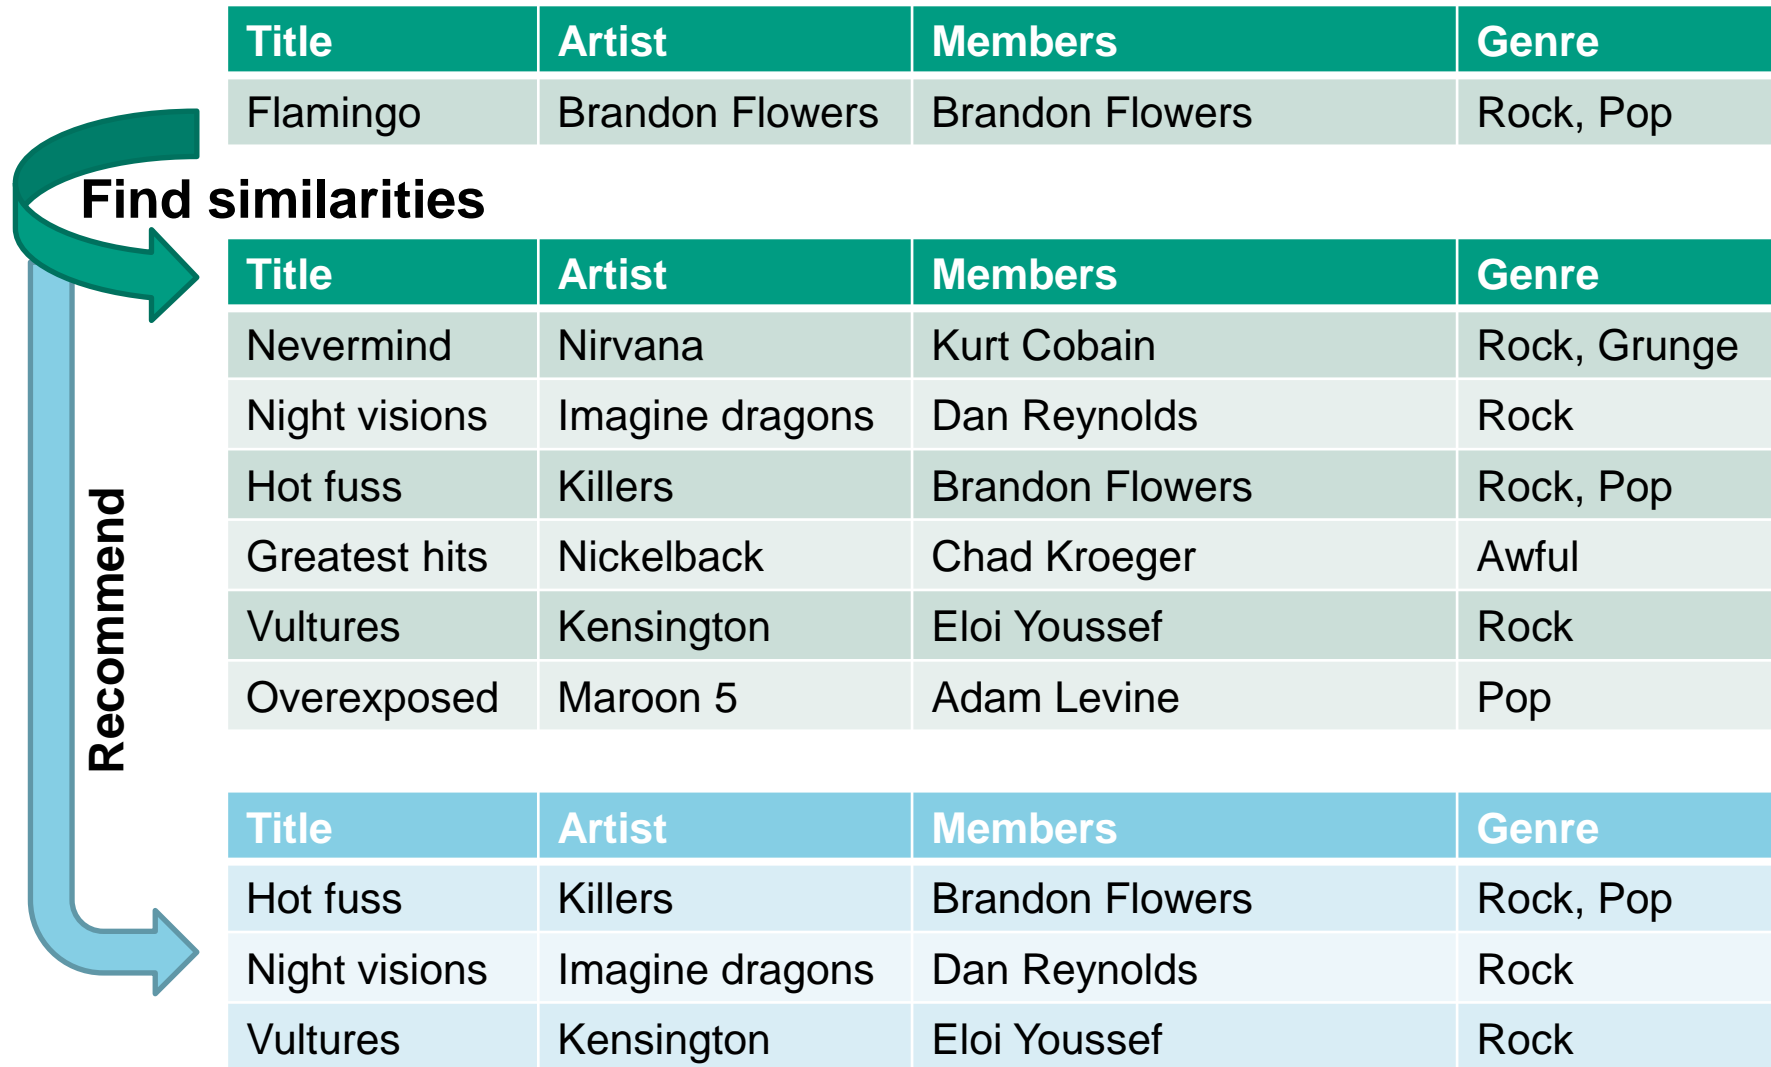
2. Create user profile



Title	Artist	Members	Genre
Flamingo	Brandon Flowers	Brandon Flowers	Rock, Pop

Content-based recommendations

3. Find similar items



Content-based recommendations

Pros

1. No need for data on other users
No **cold-start** or **sparsity** problems
2. Possible to recommend new items
No **first-rater** problem
3. Insight in the recommendation
You can understand why the item has been recommended by listing its features

Content-based recommendations

Cons


1. Finding the best features to represent items is hard
2. How to recommend to new users?
3. Overspecialization
You will never get a recommendation of items outside of your user profile



Content-based recommendations

1. How to **extract features** from CDs, books, films, ...?
2. How to calculate **similarity** between items?
3. How to validate your recommendations? (part of lecture 3)

Text learning



	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	1	1	1							
Doc 2	1		1	1	1	1				
Doc 3					1	1	1	2	1	1

Bag-of-words model

1. Get all words from a text
2. Insert words in **dictionary** (hash map)
3. Create a **sample** for each sentence
4. Each element in the vector represents a word (index from the dictionary)
5. For each word, the **number of occurrences** is stored in the vector

	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	1	1	1							
Doc 2	1		1	1	1	1				
Doc 3					1	1	1	2	1	1

Bag-of-words model Training

John likes to watch movies. Mary likes movies too.

John also likes to watch football games.

Create Feature vectors

[1, 2, 1, 1, 2, 0, 0, 0, 1, 1]

[1, 1, 1, 1, 0, 1, 1, 1, 0, 0]

Feature vectors

The word "likes"
occurs once in the 2nd
sentence

Create Dictionary



Word	Position
john	0
likes	1
to	2
watch	3
movies	4
also	5
football	6
games	7
mary	8
too	9

Dictionary

Bag-of-words model

Questions

- Does the word order matter?
- Do long phrases give different input vectors?
- Can we handle complex phrases?



Bag-of-words model

Scikit learn

Bag-of-words model a.k.a. CountVectorizer
(from sklearn.feature_extraction.text)

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
string1 = "hi Jan can you send me the presentation of week six sincerely Evert"
string2 = "hi Evert please find attached the presentation file best Jan"

# Add to list
email_list = [string1, string2]

# Create the dictionary
vectorizer.fit(email_list)
# Create feature vectors
bag_of_words = vectorizer.transform(email_list)

# Tuple: (document number, word number) => number of occurrences
print(bag_of_words)

# Print the feature number
print(vectorizer.vocabulary_.get("please"))
```

NumPy

Sparse vs dense

A **sparse matrix** is a matrix that contains mainly the value 0

For text learning (and working with ratings), it is often useful to choose a different matrix representation

NumPy has support for sparse matrices (csr_matrix)

Most operations can be performed on normal matrices and sparse matrices

Read the following article for more information:

<https://machinelearningmastery.com/sparse-matrices-for-machine-learning/>



Bag-of-words model

How to deal with...

Not all words are equal

What to do with “low-information” words?

Stop words = low information word that occurs frequently

Examples: and, the, in, for you, will, have, be

Stop words

Use NLTK (Natural Language Toolkit) package

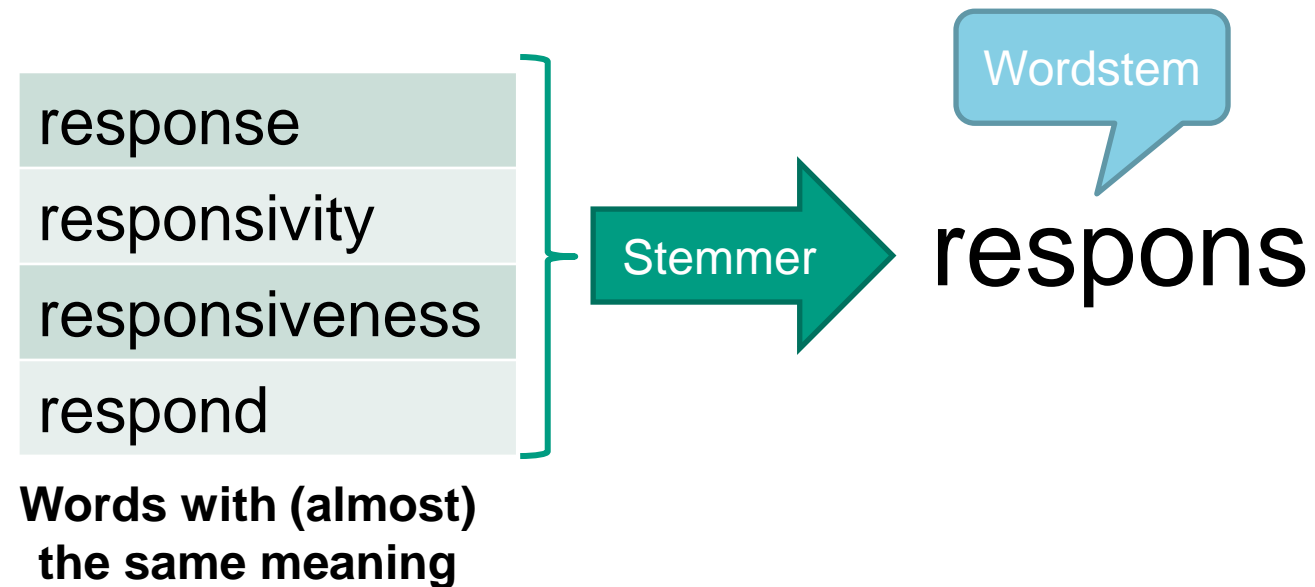
Can be used to remove stopwords from sentences

```
from nltk.corpus import stopwords
sw = stopwords.words("english")
print(sw)
```

i	is	in	now
me	are	out	d
my	was	on	ll
myself	were	off	m
we	be	over	o
our	been	under	re
ours	being	again	ve
ourselves	have	further	y
you	has	then	ain
you're	had	once	aren
you've	having	here	aren't
you'll	do	there	couldn
you'd	does	when	couldn't
your	did	where	didn
yours	doing	why	didn't
yourself	a	how	doesn
yourselves	an	all	doesn't
he	the	any	hadn
him	and	both	hadn't
his	but	each	hasn
himself	if	few	hasn't
she	or	more	haven
she's	because	most	haven't
her	as	other	isn
hers	until	some	isn't
herself	while	such	ma
it	of	no	mightn
it's	at	nor	mightn't
its	by	not	mustn
itself	for	only	mustn't
they	with	own	needn
them	about	same	needn't
their	against	so	shan
theirs	between	than	shan't
themselves	into	too	shouldn
what	through	very	shouldn't
which	during	s	wasn
who	before	t	wasn't
whom	after	can	weren
this	above	will	weren't
that	below	just	won
that'll	to	don	won't
these	from	don't	wouldn
those	up	should	wouldn't
am	down	should've	

Stemming

Not all unique words are different... Maybe we can reduce the number of possibilities?



Stemming

Python example

Use NLTK (Natural Language ToolKit) package

```
from nltk.stem.snowball import SnowballStemmer
stemmer = SnowballStemmer("english")
print(stemmer.stem("responsiveness")) # respons
print(stemmer.stem("unresponsive"))   # unrespons
```

tf-idf representation

Some words in a document are more important than others

tf-idf exposes this information

Bag of words

	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	1	1	1							
Doc 2	1		1	1	1	1				
Doc 3					1	1	1	2	1	1

tf-idf

	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	0.18	0.48	0.18							
Doc 2	0.18		0.18	0.48	0.18	0.18				
Doc 3					0.18	0.18	0.48	0.95	0.48	0.48

tf-idf representation

Reflects how important a word is to a document in a collection or corpus

tf = Term frequency

Like bag-of-words, number of occurrences of a word

idf = inverse document frequency

Weighting by how often a word occurs in the corpus

$$tf-idf(t) = tf(t) \times idf(t)$$

t is a given word

Idea: rare words score higher

tf-idf representation

tf = term frequency

Like bag-of-words, number of occurrences of a word

$$tf(t) = \frac{\text{number of times term } t \text{ appears in a document}}{\text{total number of terms in the document}}$$

John also likes to watch football games.

$$tf('also') = \frac{1}{7}$$

tf-idf representation

idf = inverse document frequency

Weighting by how often a word occurs in the corpus

$$idf(t) = \log\left(\frac{\text{total number of documents}}{\text{number of documents containing } t}\right)$$

John likes to watch movies. Mary likes movies too.

John also likes to watch football games.

$$idf('also') = \log\left(\frac{2}{1}\right) = 0.69$$

tf-idf representation

$$tf-idf(t) = tf(t) \times idf(t)$$

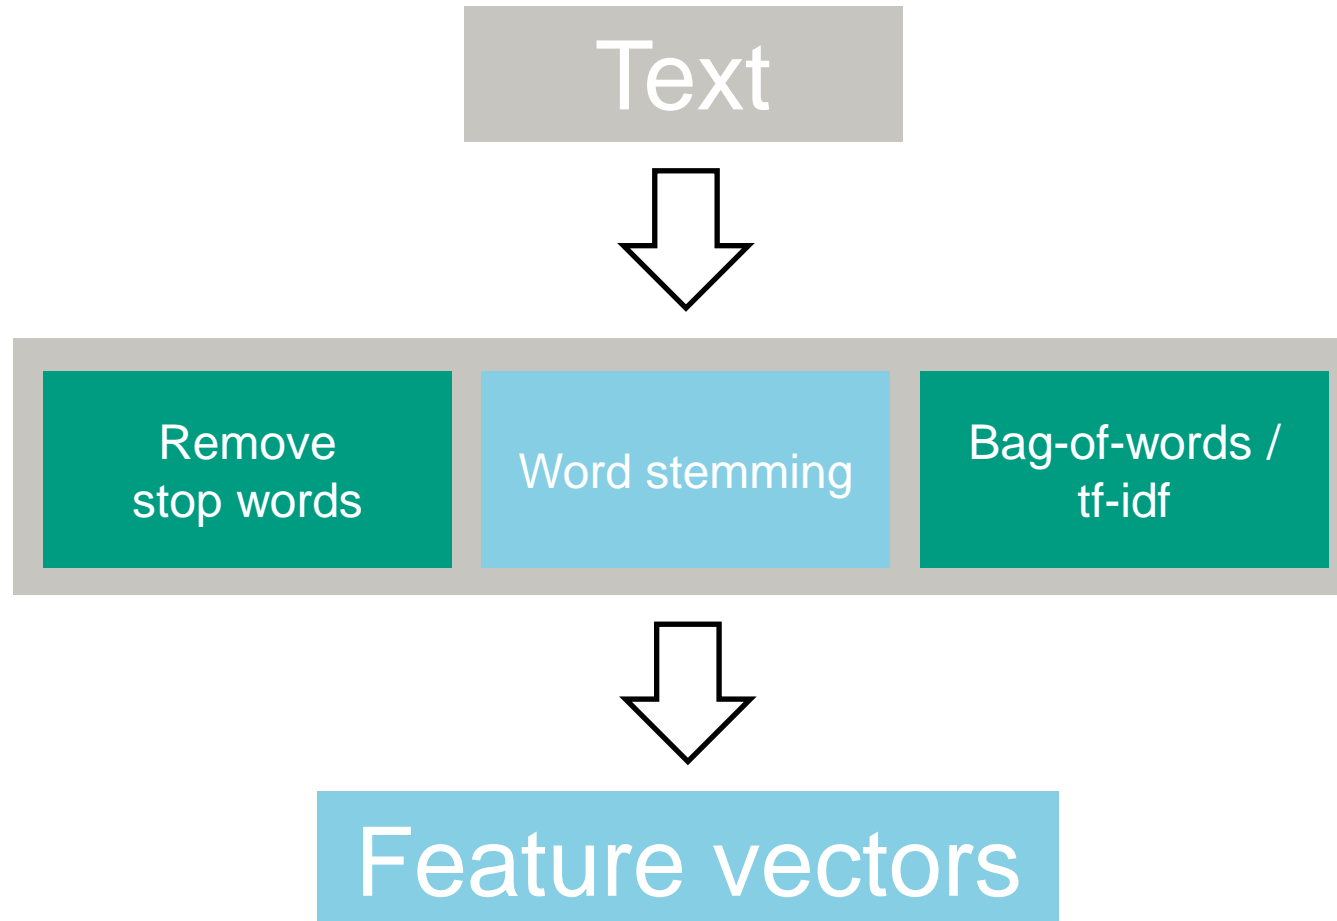
t is a given word

$$tf('also') = \frac{1}{7}$$

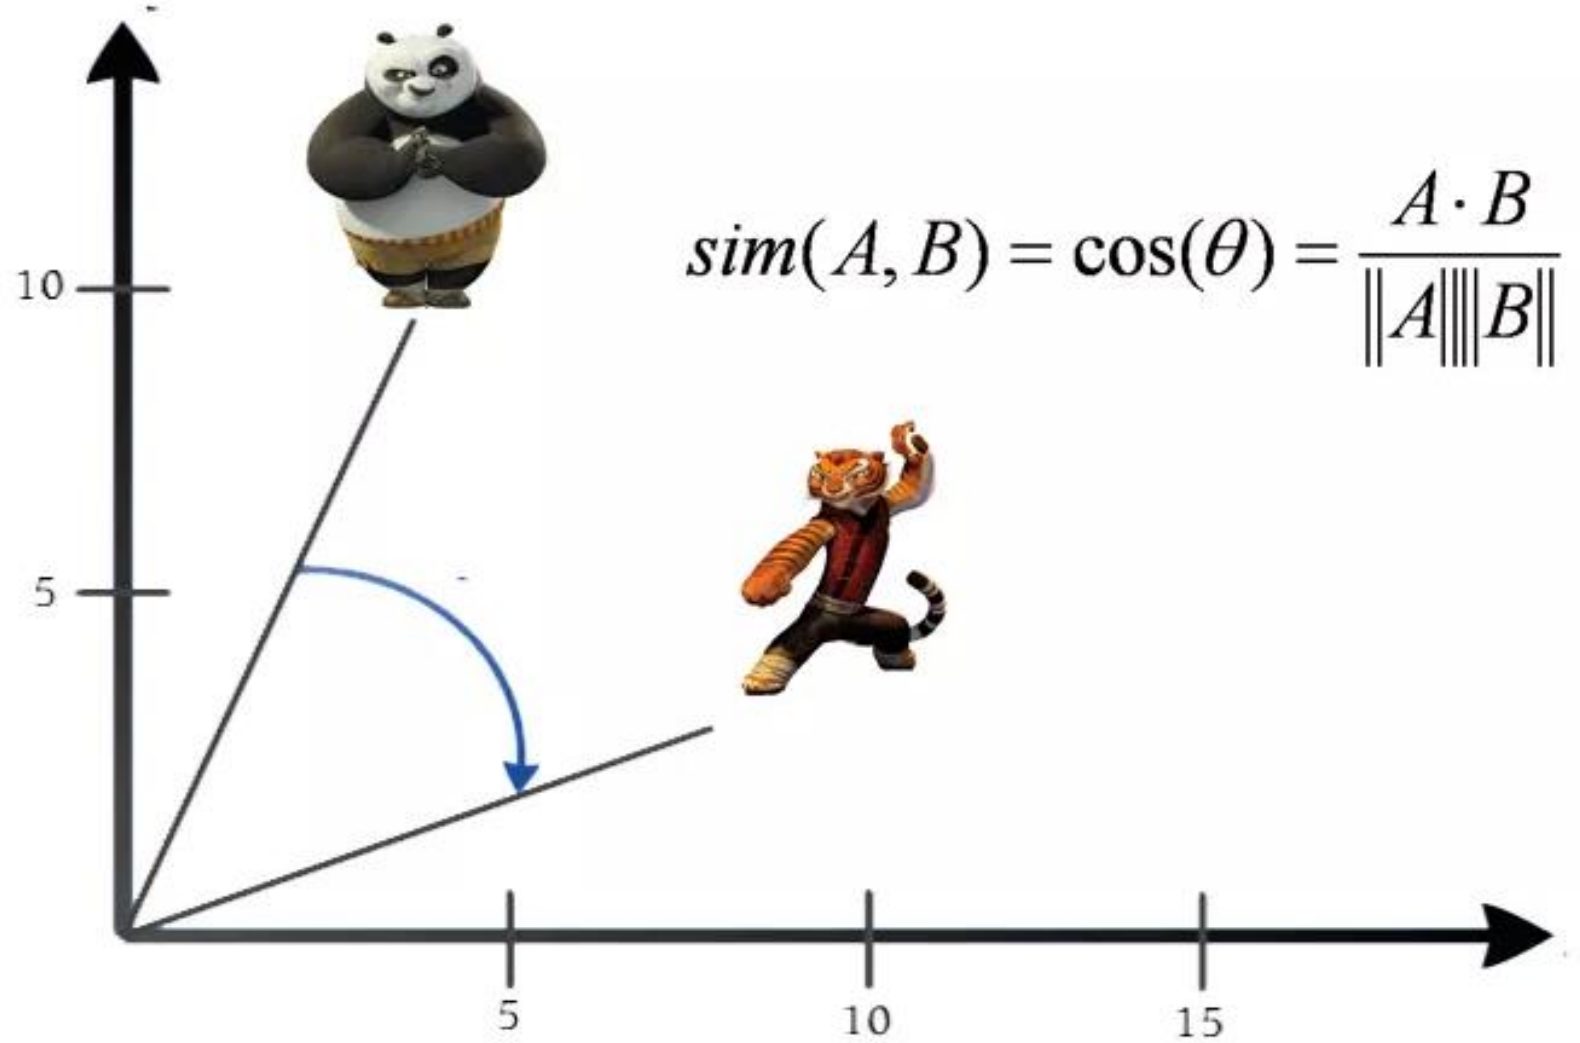
$$idf('also') = \log\left(\frac{2}{1}\right) = 0.69$$

$$tf-idf('also') = \frac{1}{7} \times 0.69 = 0.0986$$

Text learning summarized



Calculating similarities




Cosine similarity

The metric that is used most often in content-based recommendation is cosine similarity.


In short: calculate the angle between vectors. Small angles are similar items.

Artist	Genre
AC/DC	Rock
Jay Z	Rap
Linkin Park	Rock, Rap
50 Cent	Rap



Artist	Rock	Rap
AC/DC	1	0
Jay Z	0	1
Linkin Park	1	1
50 Cent	0	1


Cosine similarity



Calculating the
cosine of angle
 $\theta=45^\circ$

$$\cos(45) = 0.7071\dots$$

Cosine similarity



Calculating the
cosine of angle $\theta=0^\circ$

$$\cos(0) = 1$$

Cosine similarity

Formula

Cosine similarity = 1: items are similar

Cosine similarity = 0: items are not similar at all

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

Example similarity(Linkin Park, Jay Z)

Let vector A be an item with values [1,1]

Let vector B be an item with values [1,0]

$$\text{similarity}(A, B) = \frac{(1 * 1) + (1 * 0)}{\sqrt{1^2 + 1^2} * \sqrt{1^2 + 0^2}} = \frac{1}{\sqrt{2} * \sqrt{1}} = 0.7071$$

Cosine similarity

Formula

Example similarity(Jay Z, 50 cent)

Let vector A be an item with values [1,0]

Let vector B be an item with values [1,0]

$$\text{similarity}(A, B) = \frac{(1 * 1) + (0 * 0)}{\sqrt{1^2 + 0^2} * \sqrt{1^2 + 0^2}} = \frac{1}{1 * 1} = 1$$

Example similarity(AC-DC, Jay Z)

Let vector A be an item with values [0,1]

Let vector B be an item with values [1,0]

$$\text{similarity}(A, B) = \frac{(0 * 1) + (1 * 0)}{\sqrt{0^2 + 1^2} * \sqrt{1^2 + 0^2}} = \frac{0}{1 * 1} = 0$$

Calculating similarities



Finding the k-best matches?


Sounds like a machine learning algorithm that we have already seen...

K-nearest neighbours

We **DON'T** want the classifier variant!

Use **cosine similarity** as metric (instead of euclidian distance)

Finding the k-best matches?



Questions



Assignments

Enroll in a group!

Assignment 1: recommender systems

Many websites give users the possibility to rate items nowadays. Companies such as Amazon, Netflix, YouTube, IMDB and Bol.com use this information to recommend similar items to their users. The MovieLens dataset is a free dataset with a collection of movie ratings.

In this assignment you will build two recommendation systems, using the following techniques: content-based and collaborative filtering.

Pro-tip: finish this assignment in week 4

References

- <http://datameetsmedia.com/bag-of-words-tf-idf-explained/>
- <http://dataaspirant.com/2015/04/11/five-most-popular-similarity-measures-implementation-in-python>
- ACM Building Recommender Systems with Machine Learning and AI
- Mining Massive Datasets (mmds.org)
- Udacity Intro to Machine Learning